Streaming approaches to error detection and trimming in the analysis of short sequencing reads

TBD

December 13, 2014

1 Introduction

K-mer spectral analysis is a powerful approach to error detection and correction in shotgun sequencing data that uses k-mer abundances to determine likely errors (cite Pevzner, Quake, khmer-counting paper). Approaches derived from spectral analysis can be very accurate: Zhang et al. (2014) show that spectral analysis is considerably more effective at finding errors than quality-based approaches (cite). However, spectral analysis is also very compute intensive; most implementations must count k-mers across entire sequencing data sets, which can be memory- or I/O-intensive for large data sets.

Streaming algorithms can offer improved algorithmic and computational efficiency in the analysis of large data sets (cite). Streaming algorithms typically examine the data only once, and scale in memory usage sublinearly with respect to the input data. Streaming algorithms have not been applied to k-mer spectral analysis of reads, although XXX (melsted).

Brown et al. (2012) introduced a streaming algorithm for normalizing k-mer abundance spectra, termed "digital normalization" (abbreviated as "diginorm"). This procedure estimates the k-mer coverage of each read in an online algorithm, by calculating the median k-mer abundance of the read given all previous reads; reads above a certain estimated coverage are set aside and their k-mers are not tracked. This algorithm is both online and *streaming* because it only collects k-mers in reads with a low estimated coverage; for the same reason, it is sublinear in memory for high coverage data sets. The net effect of diginorm is to reduce the data set size that must be considered for downstream processing, such as *de novo* assembly (cite trinity, elijah, etc.)

Here we develop a streaming algorithm for k-mer spectral analysis, based on digital normalization, that can detect and remove errors in sequencing reads. This algorithm operates in sublinear memory, and examines the data at most twice. The approach offers a general framework for making use of locus-specific graph saturation and could potentially be used for error correction, variant calling, and a streaming implementation of assembly. Moreover, it can be applied to data sets with variable coverage such as transcriptomes, metagenomes, and amplified genomic DNA.

Name	Origin	Number of reads	Description
simple genome	10,000	No repeats	
E. coli MG1655	5,000,000	Subset of NCXXXXX	
simple metagenome	2,347	Dynamic range of YYY	
simple transcriptome	568	Dynamic range of ZZZ; multiple exons	

Table 1: Data sets used

2 Results

2.1 Coverage-normalized data can be used to locate and correct errors in high-coverage shotgun sequencing data

Our initial question was whether spectral error analysis could be applied to genomic short read data after digital normalization. We tested this on a synthetic data set and an *E. coli* data set. We then compared the performance of Quake on the raw and digitally normalized counts from the *E. coli* data.

Simulated data: We first applied digital normalization to a simulated data set with known errors. We generated the synthetic data set from a simulated low-complexity genome ("simple genome"; see Methods for generation and Table for data set details). We then applied digital normalization to these synthetic reads, normalizing to a median 20-mer coverage of 20 (k=20, C=20).

The k-mer spectrum before and after digital normalization is shown in Figure 1. While the total number of k-mers decreased in the digitally normalized data set, the separation between the high count k-mers and the low-count k-mers remains clear. The key concept underlying k-mer spectral error analysis is that in a high-coverage data set, these high count k-mers will represent *correct* k-mers, while the low count k-mers are produced by errors in the reads. Simple classification methods suffice to identify and trim or correct these low-count k-mers.

We next used k-mer counts from the downsampled read set to detect errors in the original read set. The algorithm is straightforward: we look for bases at the beginning or ends of low-abundance runs of k-mers in each read. We used a "trusted k-mer" cutoff of $C_0 = 3$ as our abundance cutoff, below which we assumed k-mers were erroneous (see Methods). The results are presented in Table 2.1. Of the 531 simulated reads from the simple genome containing one or more errors, predicted errors matched the known truth exactly for 443 of them (true positives), and 466 reads were correctly predicted to contain no errors (true negatives). 0 reads were falsely predicted to have no errors (false negatives). The errors in 88 reads were miscalled – while the reads each had one or more errors, the positions were not correctly called – and three reads were incorrectly predicted to contain errors, leading to a total of 91 false positives. Using the above definitions, we calculated the prediction sensitivity to be 100%

Figure 1: **K-mer spectrum of a simple artificial data set, before and after digital normalization. The peaks at 1 represents erroneous k-mers resulting from (simulated) error; the peaks centered at 80 (raw) and 20 (diginorm) represent k-mers truly present in the genome which are present in many reads.**

and the prediction specificity to be 83.0%.

When we applied spectral error detection to the unnormalized reads, we saw similar results: 440 TP, 455 TN, 105 FP, and 0 FN, for a sensitivity of 100% and a specificity of 80.7% (Table 2.1). The only difference in parameters for this analysis was that we used a cutoff of $C_0 = 10$, to account for the shifted k-mer spectrum from diginorm.

***E. coli* reads:** We next applied digital normalization and k-mer spectral error detection to an Illumina data set from *E. coli* MG1655 (cite Chitsaz). In real reads, we do not know the location of errors; to calculate them, we mapped 5m untrimmed reads to the known *E. coli* MG1655 genome with bowtie1 (cite) and recorded mismatches between the reads and the genome. These mismatches were taken as errors in the reads. We found 8.0m errors in 2.2m reads, for an overall error rate of 1.60%.

We then compared the results of k-mer spectral error detection with and without digital normalization. We used the same parameters as on the simulated genome. The results are presented in Table 2.1. Using the raw counts, the sensitivities were close – using the raw counts, we achieved a sensitivity of 72.0%, versus 71.4% using the counts from the digitally normalized reads. The specificities were also comparable – 43.5% using the raw counts, and 42.9% using the digitally normalized counts.

(Simple genome)	Raw counts	Diginorm counts
Perfect detection (TP)	440	443
No errors (TN)	455	466
Miscalled errors (FP)	91	88
Mispredicted errors (FP)	14	3
Missed errors (FN)	0	0
Sensitivity	100%	100%
Specificity	80.7%	83.0%

Table 2: Spectral error detection on 1000 synthetic reads from a simulated 10kb genome, using raw and digitally normalized counts. The counts are the number of reads where all errors were detected perfectly (TP), no errors were present (TN), one or more errors were miscalled (one type of FP), errors were mistakenly called in an error-free read (the other type of FP), and errors present in a read were missed (FN).

(E. coli)	Raw counts	Diginorm counts
Distinct k-mers	XXX	YYY
Perfect detection (TP)	818,503	807,367
No errors (TN)	2,797,925	2,797,254
Miscalled errors (FP)	1,039,052	1,046,402
Mispredicted errors (FP)	25,499	26,170
Missed errors (FN)	319,021	322,807
Sensitivity	72.0%	71.4%
Specificity	43.5%	42.9%

Table 3: Spectral error detection on 5m *E. coli* reads, using raw and digitally normalized counts.

E. coli error correction with Quake: While the results above suggest that simple spectral error detection works equally well both before and after digital normalization, we wanted to see if we could successfully perform a full error *correction* after digital normalization. To evaluate this, we applied the Quake error corrector to the data. Because Quake permits the use of a different k-mer counting table from the data being corrected, we separately used the counts from the raw reads and the counts from the digitally normalized reads to correct the raw reads (see Methods for details).

The results of running Quake on the raw data using counts from the raw or digitally normalized data are shown in Table **??**. The performance was essentially the same: Quake brought the overall error rate in the data set from 1.60% (8.0m errors) to 0.005%-0.006% (23,000 - 29,000 errors).

These results demonstrate that digitally normalized data retains all of the information necessary for effective error correction with Quake, despite having many fewer distinct k-mers.

2.2   Coverage-normalized data can be used to locate errors in variable coverage shotgun sequencing data

One of the drawbacks of spectral abundance analysis is that it cannot be applied to data with uneven coverage. For example, metagenomic or transcriptomic data sets typically contain reads from both high-abundance and low-abundance molecules. This variability in coverage confounds naive spectral analysis for two reasons: first, erroneous k-mers from very high abundance regions can accumulate and increase in abundance over the threshold for trusted k-mers, thus appearing to be correct; and second, correct reads from low coverage regions yield k-mers below the trusted k-mer threshold and appear to be incorrect. In practice, therefore, error correction for metagenomic and transcriptome data uses other approaches.

If we apply diginorm to variable coverage data and use the resulting k-mer counts to detect errors in the original reads, we can address the problem of accumulated errors from high-abundance regions: the errors do not accumulate in normalized data (cite diginorm). For the purposes of error detection, we can also avoid examining reads that are low coverage, by using the median k-mer abundance to detect low-abundance reads (Figure XXX).

@algorithm?

@graph with read abundance spectrum, showing which reads we will call errors in.

Simulated data: To test this approach, we generated two more synthetic data sets, "simple metagenome" and "simple mRNAseq," which contain both high- and low-abundance species (see Table 2.1 for data set details). After generating synthetic reads with a 1% error rate and applying digital normalization to k=20/C=20, we again did spectral error detection using the normalized counts, but with the modification that only reads with a median k-mer abundance of

20 or greater were examined.

The results of running error detection on the synthetic metagenome and mRNAseq data sets are shown in Table ZZZ. For the simple metagenome data set, 2254 of 2347 reads (96.0%) met the coverage criterion. Of the 2347 reads total, the errors in 978 erroneous reads were called perfectly (TP) and 1125 of the reads with no errors were correctly called as error-free (TN). 68 reads were incorrectly determined to be error free (FN; including the reads that were too low coverage to be considered). Of the remaining 176 reads, 170 were miscalled (errors existed but were not exactly called) and 6 were incorrectly called as erroneous when they were in fact correct. We calculated the prediction sensitivity to be 93.5% and the prediction specificity to be 84.7%. For the simple mRNAseq data set, 524 of 568 reads (92.3%) met the coverage criterion, with 228 true positives, 250 true negatives, 29 false negatives (including the low coverage reads not examined), and 61 false positives, for a prediction sensitivity of 88.7% and prediction specificity of 78.9%.

@ real data / measurements - sherine, QP.

2.3   A streaming algorithm for error detection

Even with digital normalization, the spectral error detection approach outlined above is a 2-pass offline algorithm for any given data set - the first pass normalizes the read set and records the k-mer abundances, while the second pass examines the reads. A streaming approach could avoid some or all of a second pass, leading to greater efficiency.

Here we can make use of the redundancy of shotgun sequencing to avoid examining much of the data twice. Shotgun sequencing oversamples most regions – for example, for a 100x coverage genomic data set, we would expect 50% or more of the genome to be represented in more than 100 reads. This is a consequence of the Poisson-random sampling that underlies shotgun sequencing (cite... Waterman?) This oversampling provides an opportunity, however; if we regard the read data set as a stream of incoming data, some portions of the reference will more highly sampled earlier in the stream than others. For example, in mRNAseq, highly expressed transcripts will on average be highly sampled much earlier than low-expressed transcripts.

If highly sampled regions could be detected during iteration over the read data set, we could apply the same approaches used above to do error detection in a streaming fashion. Digital normalization provides such a method: by measuring coverage for each read against an online De Bruijn graph, reads that reach a specified coverage threshold can be examined for errors immediately. Those reads that had not yet reached high coverage could be set aside and re-examined later. This could result in significant runtime savings: for a genomic data set with 100x coverage, no more than 20% of the reads would need to be examined twice.

The conceptual idea is presented in Figure 2.

In Figure 3, we show diginorm-generated coverage saturation curves for both real and error-free simulated reads from *E. coli* MG1655. In both cases, we see

Figure 2: **Diagram of streaming error detection. In a first pass over the read data, reads are loaded in until the graph locus to which they belong is saturated. From that point on, reads are examined for errors and not loaded into the graph. In a second pass, only the subset of reads loaded into the graph are examined for errors.**

that by about 1m reads, the majority of reads comes from loci with an estimated sequencing depth of 20 or higher, and hence can be used for error analysis on the first pass through the data. Only those reads collected on the first pass need to be examined again.

When we apply this streaming approach to the simulated reads from the "simple genome", we find identical numbers to the full two-pass approach: 443 TP, 466 TN, 91 FP, and 0 FN, for a sensitivity of 100% and a specificity of 83.0%. However, only 320 of the 1000 reads are examined twice - hence, it is a 1.32x pass algorithm on this data set. Likewise, for the "simple metagenome" and "simple mRNAseq" data sets, we obtain nearly identical and identical results, respectively, with the streaming approach; due to differences in the order in which reads are examined, the simple metagenome fails to detect one true positive and erroneously finds errrors in three extra reads. On the metagenome, 380 of 2347 (16.2%) of the reads are examined twice, and on the mRNAseq data set, 33.1% of the reads are examined twice.

@ do we put these numbers into a table? They're basically identical. Yes, and show deltas.

7

Figure 3: **Saturation curve of a real and a simulated *E. coli* read data set. Reads are collected when they have an estimated coverage of less than 20; in the early phase (< 1m reads), almost all reads are collected, but soon the majority of reads come from loci with an estimated sequencing depth of > 20 and are rejected.**

Sample	raw read M	normalized read M	first pass N	second pass N
simple genome	11,857	5,972 (50.4%)	1000	320 (32.0%)

Table 4: Unique k-mer counts (M) for raw and normalized data sets, and read counts (N) for first- and second-pass streaming, using a k-mer size of 20 and a coverage cutoff of 20. Digital normalization reduces the total number of k-mers being tracked, while the streaming algorithm looks at the majority of reads only once.

@ apply to E. coli?

2.4   A streaming algorithm for error trimming

Once errors can be *detected* with a streaming algorithm, it is simple to extend the algorithm to *remove* errors from reads, e.g. by trimming the read at the first error. (While it is possible to split reads around errors rather than truncating them, this introduces complications in downstream read processing.)

On the "simple genome" with counts from the digitally normalized reads, this trimming approach eliminates 149 reads entirely due to a starting low-abundance k-mer, and truncates another 392 reads. Of the 100,000 bp in the simulated reads, 31,910 (31.9%) were removed by the trimming process. In exchange, trimming eliminated *all* of the errors, bringing the overall error rate

from 0.63% to 0.00%.

For the simple metagenome we used the variable abundance approach described above and only trimmed reads with estimated coverage of 20 or higher. Here, of 2347 reads containing 234,700 bp, 314 reads (13.4%) were removed and 851 reads (36.3%) were trimmed, discarding a total of 74,321 bases (31.7%). Of 1451 errors total, all but 61 were eliminated, bringing the overall per-base error rate from 0.62% to 0.04%. The simple mRNAseq data set showed similar improvement: 83 of 568 reads were removed, and 208 were trimmed, removing 19,507 of 56,800 bases (34.34%). The initial error rate was 0.65% and the final error rate was 0.07%.

Applying the streaming error trimming to the E. coli MG1655 data set used in section XXX, we trimmed 2.0m reads and removed nearly 122,000 reads entirely. Of 7.6m errors, all but 198,000 were removed, bringing the error rate from 1.60% to 0.5%. Trimming discarded 64 Mbp of the original 500 Mbp (13.0%).

@ mrna @ real metagenome

2.5 Illumina error rates and error profiles can be determined from a small sample of sequencing data

With Illumina sequencing, average and per-position error rates may vary between sequencing runs, but are typically systematic within a run (cite?). Per-position error rates are caused by fluidics etc.

We can use the approaches described above to calculate systematic error profiles for shotgun sequencing data across entire data sets, but the variable abundance approach developed above can also be applied to *subsets* of Illumina data. The essential idea is to consume reads until sufficient data has been collected to calculate error rates, and then to calculate those error rates for the new reads based on the k-mer abundances from the old reads. This can aso be done in one pass for data sets with sufficiently high coverage data: as shown above (Figure XXX), more than half of the reads will have sufficient coverage to call errors by the time 10% of the data set has been consumed.

We first simulated a set of reads from the simple genome with errors only at even positions (0, 2, 4, etc.), and called errors in these reads using the single-pass algorithm described above (see Methods for implementation details). We also separately called errors by mapping with bowtie and examining mismatches. The results are shown in Figure YYY; the difference in predicted vs actual mismatch profiles has an average of 0.04%, with a variance of 0.0002% across positions. The predicted per-base error rate is 0.57%, while the true error rate is 0.65%; this underestimate may be due to ignoring reads with multiple errors in them that present as too low-coverage to assess.

@@show figure with shaded color indicating difference

We next applied to ecoli data.@@

@@error rates? apply regression?

3 Discussion

3.1 Digital normalization enables k-mer spectral error detection in variable abundance shotgun data

Reference-free error detection and correction in shotgun data from metagenomic, transcriptomic, and amplified samples typically requires specialized approaches (cite, discuss). Here we demonstrate that after digital normalization, k-mer abundance spectral approaches can be applied to reads with high estimated coverage. This should generalize to any approaches that rely on k-mer abundance, including error detection, trimming, correction.

One significant disadvantage of the variable abundance approach is that low-coverage reads cannot be trimmed because we do not know whether they are from low-abundance molecules or are highly erroneous. This must be taken into account for downstream analysis; for example, while assemblers must already ignore or correct these erroneous reads, quantitation approaches using e.g. mapping should be not be applied to the trimmed data.

Another disadvantage of the variable abundance approach is that we cannot distinguish between low abundance and highly erroneous reads because the median k-mer abundance is an underestimate of coverage for highly erroneous reads. (In k-mer spectral error analysis of genomes, low abundance reads can simply be assumed to be highly erroneous.)

(@This is where applying Quake to diginorm reads would be a useful demonstration.)

3.2 A streaming algorithm for read analysis

We introduce a few-pass algorithm ($\leq 2N$) for detecting and analyzing reads that belong to high-coverage regions of the De Bruijn graph. As with digital normalization, the algorithm is sublinear in its memory use because only k-mer counts in the retained reads are tracked. The algorithm looks at every read once, and some reads twice; the exact number is dictated by biological features of the data set as well as the error rate, but for several real data sets is below 30% (Table 2.3).

This algorithm can potentially be applied to any shotgun sequencing data set, including genomic, transcriptomic, metagenomic, and whole-genome amplified data.

If better methods for read coverage estimation come along, they can be used in place of median k-mer abundance.

Potentially opens up streaming error correction and variant calling.

3.3 Empirical error profiles can be calculated on subsets of Illumina shotgun data

XXXX

Useful for cores etc: quick evaluation of sequencing quality, regardless of origin.

Repeats could be pulled off first.

3.4   Concluding thoughts

Here we have demonstrated that a general streaming algorithm based on detection of graph-local saturation can be used to detect and remove errors in high-coverage shotgun sequencing data. The algorithm is few-pass, in that some fraction of the data must be examined twice, although the precise fraction depends on biological features of the data set. In addition, because the algorithm relies on abundance-normalized sequencing data, it can be applied to metagenomic and transcriptomic data as well as genomic. And, finally, the approach may enable some sublinear-time analyses as well.

From a pragmatic perspective, the error trimming approaches described here should only be used on metagenomes and transcriptomes prior to assembly, and should not be used prior to mapping. Because they ignore low-abundance reads, those reads will not be trimmed, and may map poorly, in which case downstream analyses such as counting may be biased against low-abundance reads.

Time- and memory-efficient error profile calculation is practically useful for large sequencing cores.

Dealing with variants is outside the scope of this work, but 50/50 variants should not be trimmed. What about repeats?

Streaming methods for error correction and variant calling, more generally.

Theory needed for thresholding, but, in theory, thresholds should be static. (We can demonstrate that there is little variation.) Interesting to examine theory, real parameters for data sets... for the future.